

Case: Sensor Data Modeling

Grupo 1:

Marcelo Barbugli

Gisele Siqueira

Diego Moura

Roberto Eyama

Ricardo Geroto

Matheus Higa

Atividade feita utilizando as plataformas DataStax e AWS Cloud9:

Evidências de:

- Criar as tabelas
- Popular dados
- Realizar consultas

(extra) Step1: Criando um banco de dados no Astra (datastax)

```
get Cloud provider
[INFO] Database 'data-modeling' and keyspace 'sensor_data' are being created.
[INFO] Database 'data-modeling' has status 'PENDING' waiting to be 'ACTIVE' ...
[INFO] Database 'data-modeling' has status 'ACTIVE' (took 111861 millis)
[OK] Database 'data-modeling' is ready.
gitpod /workspace/data-modeling-sensor-data (main) $ astra db list
```

Name	id	Regions	Cloud	V	Status
--wait	f5efef7c-30bb-494d-95ba-f6573f558ca3	us-east1	gcp		ACTIVE
data-modeling	e2c09b8d-337b-4bfb-8ff9-13e012afd02b	us-east1	gcp		ACTIVE

```
gitpod /workspace/data-modeling-sensor-data (main) $ astra db get data-modeling
```

Attribute	Value
Name	data-modeling
id	e2c09b8d-337b-4bfb-8ff9-13e012afd02b
Cloud	GCP
Regions	us-east1
Status	ACTIVE
Vector	Disabled
Default Keyspace	sensor_data
Creation Time	2024-06-06T00:38:31Z
Keyspaces	[0] sensor_data
Regions	[0] us-east1

```
gitpod /workspace/data-modeling-sensor-data (main) $ astra db cqlsh data-modeling -k sensor_data
[INFO] Cqlsh is starting, please wait for connection establishment...
Connected to cndb at 127.0.0.1:9042.
[cqlsh 6.8.0 | Cassandra 4.0.0.6816 | CQL spec 3.4.5 | Native protocol v4]
Use HELP for help.
token@cqlsh:sensor_data> 7w
```

Astra Dashboard – Serverless Database Service (SaaS)

Db: data-modeling

Keyspace: sensor_data

data-modeling Active

Database ID: e2c09b8d-337b-4bfb-8ff9-13e012afd02b

Overview Health Connect CQL Console CDC Settings

Usage for Current Billing Period ⓘ

Read Requests	Write Requests	Storage Consumed ⓘ	Data Transfer
0	0	10.7 KB	103.43 KB

Regions Unlock Multi-Region

Paid plans allow you to choose multiple regions.

Provider	Area	Region	Region Name	Datacenter ID	Region Availability
Google Cloud	North America	us-east1	Moncks Corner, South...	e2c09b8d...2b-1	Online

Keyspaces Add Keyspace

Learn more about [keyspaces](#) and how to use them.

Keyspace
sensor_data

Step1: Criando banco de dados Cassandra via CQL shell, criando a KEYSPACE 'sensor_data' e passando a chamada para iniciar a criação das tabelas dentro deste *keyspace*:

DataStax:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 4 EXPOSED PORTS AZURE COMMENTS
./cassandra gitpod /workspace/data-modeling-sensor-data (main) $ ./cassandra
Starting a Cassandra cluster ... DONE!
Cassandra successfully started.
gitpod /workspace/data-modeling-sensor-data (main) $ cqlsh
WARNING: cqlsh was built against 4.1.4, but this server is 4.0.13. All features may not work!
Connected to Cassandra Cluster at 127.0.0.1:9042
[cqlsh 6.1.0 | Cassandra 4.0.13 | CQL spec 3.4.5 | Native protocol v5]
Use HELP for help.
cqlsh> CREATE KEYSPACE sensor_data
'class': 'NetworkTopologyStrategy',
  'DC-Housto ... WITH replication = {
    ... 'class': 'NetworkTopologyStrategy',
    ... 'DC-Houston': 1 };
cqlsh> USE sensor_data;
cqlsh:sensor_data>
```

AWS:

```
cqlsh> DROP KEYSPACE IF EXISTS killrvideo;

cqlsh> // CREATE KEYSPACE killrvideo WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 3 };
cqlsh> CREATE KEYSPACE sensor_data WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 1 };
cqlsh>
cqlsh>
cqlsh> USE sensor_data;
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS sensor_data (
```

Step 2: Criando as tabelas `networks`, `sensors_by_network`, `temperatures_by_network`, e `temperatures_by_sensor` e executando `DESCRIBE` para confirmar a criação. A PRIMARY KEY é composta pela Partition Key + Clustering key. Exemplo da tabela `networks` é Partition Key = `bucket`; Clustering key = `name` e assim formando a PRIMARY KEY.

DataStax:

```
use HELP FOR help.
cqlsh> CREATE KEYSPACE sensor_data
'class': 'NetworkTopologyStrategy',
'DC-Housto ... WITH replication = {
... 'class': 'NetworkTopologyStrategy',
... 'DC-Houston': 1 };
cqlsh> USE sensor_data;
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS networks (
... bucket TEXT,
... name TEXT,
... description TEXT,
... region TEXT,
... num_sensors INT,
... PRIMARY KEY ((bucket),name)
... );
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS temperatures_by_network (
... network TEXT,
... week DATE,
... date_hour TIMESTAMP,
... sensor TEXT,
... avg_temperature FLOAT,
... latitude DECIMAL,
... longitude DECIMAL,
... PRIMARY KEY ((network,week),date_hour,sensor)
... ) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS sensors_by_network (
... network TEXT,
... sensor TEXT,
... latitude DECIMAL,
... longitude DECIMAL,
... characteristics MAP<TEXT,TEXT>,
... PRIMARY KEY ((network),sensor)
... );
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS temperatures_by_sensor (
... sensor TEXT,
... date DATE,
... timestamp TIMESTAMP,
... value FLOAT,
... PRIMARY KEY ((sensor,date),timestamp)
... ) WITH CLUSTERING ORDER BY (timestamp DESC);
cqlsh:sensor_data> DESCRIBE TABLES;

networks  sensors_by_network  temperatures_by_network  temperatures_by_sensor
```

AWS:

```

cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS networks (
...     bucket TEXT,
...     name TEXT,
...     description TEXT,
...     region TEXT,
...     num_sensors INT,
...     PRIMARY KEY ((bucket),name)
... );
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS temperatures_by_network (
...     network TEXT,
...     week DATE,
...     date_hour TIMESTAMP,
...     sensor TEXT,
...     avg_temperature FLOAT,
...     latitude DECIMAL,
...     longitude DECIMAL,
...     PRIMARY KEY ((network,week),date_hour,sensor)
... ) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS sensors_by_network (
...     network TEXT,
...     sensor TEXT,
...     latitude DECIMAL,
...     longitude DECIMAL,
...     characteristics MAP<TEXT,TEXT>,
...     PRIMARY KEY ((network),sensor)
... );
cqlsh:sensor_data> CREATE TABLE IF NOT EXISTS temperatures_by_sensor (
...     sensor TEXT,
...     date DATE,
...     timestamp TIMESTAMP,
...     value FLOAT,
...     PRIMARY KEY ((sensor,date),timestamp)
... ) WITH CLUSTERING ORDER BY (timestamp DESC);CREATE TABLE IF NOT EXISTS temperatures_by_sensor (
...     sensor TEXT,
...     date DATE,
...     timestamp TIMESTAMP,
...     value FLOAT,
...     PRIMARY KEY ((sensor,date),timestamp)
... ) WITH CLUSTERING ORDER BY (timestamp DESC);CREATE TABLE IF NOT EXISTS temperatures_by_sensor (
...     sensor TEXT,
...     date DATE,
...     timestamp TIMESTAMP,
...     value FLOAT,
...     PRIMARY KEY ((sensor,date),timestamp)
... ) WITH CLUSTERING ORDER BY (timestamp DESC);
cqlsh:sensor_data>
cqlsh:sensor_data> DESCRIBE TABLES;

networks  sensors_by_network  temperatures_by_network  temperatures_by_sensor

```

STEP3: Execução de script para realizar INSERT nas 4 tabelas.

DataStax:

```

cqlsh:sensor_data> source '/workspace/data-modeling-sensor-data/assets/sensor_data.cql'
WARNING: cqlsh was built against 4.1.4, but this server is 4.0.13. All features may not work!
1 table created: SELECT * FROM networks;

```

AWS:

```

cqlsh:sensor_data> -- Populate table networks:
cqlsh:sensor_data> -----
cqlsh:sensor_data> INSERT INTO networks
... (bucket,name,description,region,num_sensors)
... VALUES ('all','forest-net',
...          'forest fire detection network',
...          'south',3);
cqlsh:sensor_data> INSERT INTO networks
... (bucket,name,description,region,num_sensors)
... VALUES ('all','volcano-net',
...          'volcano monitoring network',
...          'north',2);
cqlsh:sensor_data> INSERT INTO tcqlsh:sensor_data>
cqlsh:sensor_data>
cqlsh:sensor_data> -- Populate table sensors_by_network:
cqlsh:sensor_data> -----
cqlsh:sensor_data> INSERT INTO sensors_by_network
itude)
VALUES ('forest-net', ... (network,sensor,latitude,longitude,characteristics)
... VALUES ('forest-net','s1001',30.526503,-95.582815,
...          {'accuracy':'medium','sensitivity':'high'});
cqlsh:sensor_data> INSERT INTO sensors_by_network
... (network,sensor,latitude,longitude,characteristics)
... VALUES ('forest-net','s1002',30.518650,-95.583585,
...          {'accuracy':'medium','sensitivity':'high'});
cqlsh:sensor_data> INSERT INTO sensors_by_network
tude)
VAL ... (network,sensor,latitude,longitude,characteristics)
... VALUES ('forest-net','s1003',30.515056,-95.556225,
...          {'accuracy':'medium','sensitivity':'high'});

-----cqlsh:sensor_data> INSERT INTO sensors_by_network
... (network,sensor,latitude,longitude,characteristics)
... VALUES ('volcano-net','s2001',44.460321,-110.828151,
...          {'accuracy':'high','sensitivity':'medium'});
cqlsh:sensor_data> INSERT INTO sensors_by_network

```

Realizando SELECT * na tabela networks: Partition Key = bucket; Clustering key = name; outros atributos = description, num_sensors, region.

DataStax:

```

cqlsh:sensor_data> SELECT * FROM networks;

```

bucket	name	description	num_sensors	region
all	forest-net	forest fire detection network	3	south
all	volcano-net	volcano monitoring network	2	north

AWS:

```
cqlsh:sensor_data> SELECT * FROM networks;
```

bucket	name	description	num_sensors	region
all	forest-net	forest fire detection network	3	south
all	volcano-net	volcano monitoring network	2	north

(2 rows)

Realizando SELECT, passando duas colunas que são as Partition Key (network e week), duas colunas que são as Clustering Key (date_hour e sensor), onde juntas formam a PRIMARY KEY. Outro atributo é a coluna avg_temperature.

DataStax:

```
cqlsh:sensor_data> SELECT network, week, date_hour,
... sensor, avg_temperature
... FROM temperatures_by_network;
```

network	week	date_hour	sensor	avg_temperature
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1001	97.5
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1002	100
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1003	98.5
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1001	79.5
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1002	81
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1003	80.5
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1001	106.5
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1002	109
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1003	1372
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1001	90.5
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1002	90
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1003	90.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1001	98.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1002	99.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1003	101.5
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1001	80.5
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1002	82
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1003	82.5

(18 rows)

AWS:

```
cqlsh:sensor_data> SELECT network, week, date_hour,
...      sensor, avg_temperature
... FROM temperatures_by_network;
```

network	week	date_hour	sensor	avg_temperature
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1001	97.5
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1002	100
forest-net	2020-06-28	2020-07-04 12:00:00.000000+0000	s1003	98.5
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1001	79.5
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1002	81
forest-net	2020-06-28	2020-07-04 00:00:00.000000+0000	s1003	80.5
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1001	106.5
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1002	109
forest-net	2020-07-05	2020-07-06 12:00:00.000000+0000	s1003	1372
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1001	90.5
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1002	90
forest-net	2020-07-05	2020-07-06 00:00:00.000000+0000	s1003	90.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1001	98.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1002	99.5
forest-net	2020-07-05	2020-07-05 12:00:00.000000+0000	s1003	101.5
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1001	80.5
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1002	82
forest-net	2020-07-05	2020-07-05 00:00:00.000000+0000	s1003	82.5

(18 rows)

Realizando SELECT * da tabela sensors_by_network. Partition Key = network; Clustering key = sensor; outros atributos = characteristics, latitude, longitude.

DataStax:

```
cqlsh:sensor_data> SELECT * FROM sensors_by_network;
```

network	sensor	characteristics	latitude	longitude
forest-net	s1001	{'accuracy': 'medium', 'sensitivity': 'high'}	30.526503	-95.582815
forest-net	s1002	{'accuracy': 'medium', 'sensitivity': 'high'}	30.518650	-95.583585
forest-net	s1003	{'accuracy': 'medium', 'sensitivity': 'high'}	30.515056	-95.556225
volcano-net	s2001	{'accuracy': 'high', 'sensitivity': 'medium'}	44.460321	-110.828151
volcano-net	s2002	{'accuracy': 'high', 'sensitivity': 'medium'}	44.463195	-110.830124

(5 rows)

AWS:

```
cqlsh:sensor_data> SELECT * FROM sensors_by_network;
```

network	sensor	characteristics	latitude	longitude
forest-net	s1001	{'accuracy': 'medium', 'sensitivity': 'high'}	30.526503	-95.582815
forest-net	s1002	{'accuracy': 'medium', 'sensitivity': 'high'}	30.518650	-95.583585
forest-net	s1003	{'accuracy': 'medium', 'sensitivity': 'high'}	30.515056	-95.556225
volcano-net	s2001	{'accuracy': 'high', 'sensitivity': 'medium'}	44.460321	-110.828151
volcano-net	s2002	{'accuracy': 'high', 'sensitivity': 'medium'}	44.463195	-110.830124

(5 rows)

Realizando SELECT * da tabela temperature_by_sensor. Partition Key = sensor e date; Clustering key = timestamp; outros atributos = value.

DataStax:

```
cqlsh:sensor_data> SELECT * FROM temperatures_by_sensor;
```

sensor	date	timestamp	value
s1001	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1001	2020-07-04	2020-07-04 12:00:01.000000+0000	97
s1001	2020-07-04	2020-07-04 00:59:59.000000+0000	79
s1001	2020-07-04	2020-07-04 00:00:01.000000+0000	80
s1001	2020-07-05	2020-07-05 12:59:59.000000+0000	99
s1001	2020-07-05	2020-07-05 12:00:01.000000+0000	98
s1001	2020-07-05	2020-07-05 00:59:59.000000+0000	80
s1001	2020-07-05	2020-07-05 00:00:01.000000+0000	81
s1002	2020-07-06	2020-07-06 12:59:59.000000+0000	110
s1002	2020-07-06	2020-07-06 12:00:01.000000+0000	108
s1002	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1002	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1003	2020-07-04	2020-07-04 12:00:01.000000+0000	99
s1003	2020-07-04	2020-07-04 00:59:59.000000+0000	80
s1003	2020-07-04	2020-07-04 00:00:01.000000+0000	81
s1003	2020-07-06	2020-07-06 12:59:59.000000+0000	1429
s1003	2020-07-06	2020-07-06 12:00:01.000000+0000	1315
s1003	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1003	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-05	2020-07-05 12:59:59.000000+0000	102
s1003	2020-07-05	2020-07-05 12:00:01.000000+0000	101
s1003	2020-07-05	2020-07-05 00:59:59.000000+0000	82

AWS:


```
cqlsh:sensor_data> SELECT * FROM temperatures_by_sensor;
```

sensor	date	timestamp	value
s1001	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1001	2020-07-04	2020-07-04 12:00:01.000000+0000	97
s1001	2020-07-04	2020-07-04 00:59:59.000000+0000	79
s1001	2020-07-04	2020-07-04 00:00:01.000000+0000	80
s1001	2020-07-05	2020-07-05 12:59:59.000000+0000	99
s1001	2020-07-05	2020-07-05 12:00:01.000000+0000	98
s1001	2020-07-05	2020-07-05 00:59:59.000000+0000	80
s1001	2020-07-05	2020-07-05 00:00:01.000000+0000	81
s1002	2020-07-06	2020-07-06 12:59:59.000000+0000	110
s1002	2020-07-06	2020-07-06 12:00:01.000000+0000	108
s1002	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1002	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-04	2020-07-04 12:59:59.000000+0000	98
s1003	2020-07-04	2020-07-04 12:00:01.000000+0000	99
s1003	2020-07-04	2020-07-04 00:59:59.000000+0000	80
s1003	2020-07-04	2020-07-04 00:00:01.000000+0000	81
s1003	2020-07-06	2020-07-06 12:59:59.000000+0000	1429
s1003	2020-07-06	2020-07-06 12:00:01.000000+0000	1315
s1003	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1003	2020-07-06	2020-07-06 00:00:01.000000+0000	90
s1003	2020-07-05	2020-07-05 12:59:59.000000+0000	102
s1003	2020-07-05	2020-07-05 12:00:01.000000+0000	101
s1003	2020-07-05	2020-07-05 00:59:59.000000+0000	82
s1003	2020-07-05	2020-07-05 00:00:01.000000+0000	83
s1002	2020-07-05	2020-07-05 12:59:59.000000+0000	99
s1002	2020-07-05	2020-07-05 12:00:01.000000+0000	100
s1002	2020-07-05	2020-07-05 00:59:59.000000+0000	82
s1002	2020-07-05	2020-07-05 00:00:01.000000+0000	82
s1002	2020-07-04	2020-07-04 12:59:59.000000+0000	100
s1002	2020-07-04	2020-07-04 12:00:01.000000+0000	100
s1002	2020-07-04	2020-07-04 00:59:59.000000+0000	80
s1002	2020-07-04	2020-07-04 00:00:01.000000+0000	82
s1001	2020-07-06	2020-07-06 12:59:59.000000+0000	107
s1001	2020-07-06	2020-07-06 12:00:01.000000+0000	106
s1001	2020-07-06	2020-07-06 00:59:59.000000+0000	90
s1001	2020-07-06	2020-07-06 00:00:01.000000+0000	90

(36 rows)

STEP4: Buscando todas as informações da tabela networks, ordenando pela Clustering key = name.

Realizando SELECT nas colunas name, description, region, num_sensors da tabela networks:

A importância do WHERE é porque bucket representa a Partition Key da tabela networks, sendo necessário passar na query a condição de igualdade (ou desigualdade) com valor atribuído, ordenando e garantindo unicidade quando a tabela foi criada.

DataStax:

```
cqlsh:sensor_data> SELECT name, description,
...         region, num_sensors
... FROM networks
... WHERE bucket = 'all';
```

name	description	region	num_sensors
forest-net	forest fire detection network	south	3
volcano-net	volcano monitoring network	north	2

AWS:

```
cqlsh:sensor_data> SELECT name, description,
...         region, num_sensors
... FROM networks
... WHERE bucket = 'all';
```

name	description	region	num_sensors
forest-net	forest fire detection network	south	3
volcano-net	volcano monitoring network	north	2

(2 rows)

STEP5: Encontre as temperaturas médias horárias para cada sensor na network = **forest-net** e intervalo de datas [2020-07-05, 2020-07-06] dentro da semana de 2020-07-05; order by date (desc) e hour (desc).

DataStax:

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
...         latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week = '2020-07-05'
... AND date_hour >= '2020-07-05'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003

(12 rows)

AWS:

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
... latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week = '2020-07-05'
... AND date_hour >= '2020-07-05'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003

(12 rows)

Encontre as temperaturas médias horárias para cada sensor na network = **forest-net** e intervalo de datas [2020-07-04,2020-07-06] entre as semanas de 2020-06-28 e 2020-07-05; order by date (desc) e hour (desc).

Solution1:

DataStax:

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
... latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week = '2020-06-28'
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003

(6 rows)

AWS:

```

cqlsh:sensor_data>
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
...         latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week = '2020-06-28'
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';

```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003

(6 rows)

Solution2:

DataStax:

```

cqlsh:sensor_data> SELECT date_hour, avg_temperature,
...         latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week IN ('2020-07-05', '2020-06-28')
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';

```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003

(18 rows)

AWS:

```
cqlsh:sensor_data> SELECT date_hour, avg_temperature,
... latitude, longitude, sensor
... FROM temperatures_by_network
... WHERE network = 'forest-net'
... AND week IN ('2020-07-05','2020-06-28')
... AND date_hour >= '2020-07-04'
... AND date_hour < '2020-07-07';
```

date_hour	avg_temperature	latitude	longitude	sensor
2020-07-04 12:00:00.000000+0000	97.5	30.526503	-95.582815	s1001
2020-07-04 12:00:00.000000+0000	100	30.518650	-95.583585	s1002
2020-07-04 12:00:00.000000+0000	98.5	30.515056	-95.556225	s1003
2020-07-04 00:00:00.000000+0000	79.5	30.526503	-95.582815	s1001
2020-07-04 00:00:00.000000+0000	81	30.518650	-95.583585	s1002
2020-07-04 00:00:00.000000+0000	80.5	30.515056	-95.556225	s1003
2020-07-06 12:00:00.000000+0000	106.5	30.526503	-95.582815	s1001
2020-07-06 12:00:00.000000+0000	109	30.518650	-95.583585	s1002
2020-07-06 12:00:00.000000+0000	1372	30.515056	-95.556225	s1003
2020-07-06 00:00:00.000000+0000	90.5	30.526503	-95.582815	s1001
2020-07-06 00:00:00.000000+0000	90	30.518650	-95.583585	s1002
2020-07-06 00:00:00.000000+0000	90.5	30.515056	-95.556225	s1003
2020-07-05 12:00:00.000000+0000	98.5	30.526503	-95.582815	s1001
2020-07-05 12:00:00.000000+0000	99.5	30.518650	-95.583585	s1002
2020-07-05 12:00:00.000000+0000	101.5	30.515056	-95.556225	s1003
2020-07-05 00:00:00.000000+0000	80.5	30.526503	-95.582815	s1001
2020-07-05 00:00:00.000000+0000	82	30.518650	-95.583585	s1002
2020-07-05 00:00:00.000000+0000	82.5	30.515056	-95.556225	s1003

(18 rows)

STEP6: Buscando informações de todos os sensors (Clustering Key) em network (Partition Key) onde o valor seja = 'forest-net'.

DataStax:

```
cqlsh:sensor_data> SELECT *
... FROM sensors_by_network
... WHERE network = 'forest-net';
```

network	sensor	characteristics	latitude	longitude
forest-net	s1001	{'accuracy': 'medium', 'sensitivity': 'high'}	30.526503	-95.582815
forest-net	s1002	{'accuracy': 'medium', 'sensitivity': 'high'}	30.518650	-95.583585
forest-net	s1003	{'accuracy': 'medium', 'sensitivity': 'high'}	30.515056	-95.556225

(3 rows)

AWS:

```
cqlsh:sensor_data> SELECT *
... FROM sensors_by_network
... WHERE network = 'forest-net';
```

network	sensor	characteristics	latitude	longitude
forest-net	s1001	{'accuracy': 'medium', 'sensitivity': 'high'}	30.526503	-95.582815
forest-net	s1002	{'accuracy': 'medium', 'sensitivity': 'high'}	30.518650	-95.583585
forest-net	s1003	{'accuracy': 'medium', 'sensitivity': 'high'}	30.515056	-95.556225

(3 rows)

STEP7: Encontre medições brutas para o sensor **s1003** em 2020-07-06; order by **timestamp** (desc).

DataStax:

```
cqlsh:sensor_data> SELECT timestamp, value
... FROM temperatures_by_sensor
... WHERE sensor = 's1003'
... AND date = '2020-07-06';
```

timestamp	value
2020-07-06 12:59:59.000000+0000	1429
2020-07-06 12:00:01.000000+0000	1315
2020-07-06 00:59:59.000000+0000	90
2020-07-06 00:00:01.000000+0000	90

(4 rows)

AWS:

```
cqlsh:sensor_data> SELECT timestamp, value
... FROM temperatures_by_sensor
... WHERE sensor = 's1003'
... AND date = '2020-07-06';
```

timestamp	value
2020-07-06 12:59:59.000000+0000	1429
2020-07-06 12:00:01.000000+0000	1315
2020-07-06 00:59:59.000000+0000	90
2020-07-06 00:00:01.000000+0000	90

(4 rows)

Consideração final entre DataStax e AWS Cloud9:

1. A diferença que encontramos foi ao criar imagem do Cassandra no Cloud9, onde é necessário. Já no dataStax precisamos apenas iniciar pois a imagem já está montada.
2. Para criar o KEYSpace, a sintaxe entre DataStax e AWS Cloud9 sofre alteração, pois no DataStax permite a criação da classe da réplica = 'NetworkTopologyStrategy', enquanto nossa conta na AWS Academy só permite o

'SimpleStrategy'. Outro fator é que pelo DataStax (GCP) passamos a região que nossa replicação estará alocada e pelo AWS Cloud9 podemos escolher o 'replication_factor' = 1 (ou 2, 3) região(ões). E.g.:

- a. DataStax: CREATE KEYSPACE sensor_data WITH REPLICATION = { 'class': 'NetworkTopologyStrategy', 'DC-Houston': 1 };
- b. AWS Cloud9: CREATE KEYSPACE sensor_data WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 1 };